

Deep Reinforcement Learning Project 2 : Continuous Control

Rachel Schlossman

September 8, 2021

1 Learning Algorithm

The learning algorithm is an implementation of the Deep Deterministic Policy Gradient algorithm as described in [2]. As the paper describes, the algorithm leverages (1) an actor network to approximate actor function, $\mu(s|\theta^\mu)$, and a critic network to approximate the action-value function, $Q(s, a|\theta^Q)$, using the Bellman Equation (similar to Q-learning). The variable s is the state, a is the action, and θ^μ and θ^Q are the actor and critic network weights, respectively. The algorithm is suited only for continuous action spaces.

1.1 Hyperparameters

The following hyperparameters were used in the DDPG implementation:

- replay buffer size : 1e6
- minibatch size: 128
- discount factor: 0.99
- for soft update of target parameters, τ : 1e-3
- learning rate of the actor: 1e-4
- learning rate of the critic: 1e-4
- L2 weight decay: 0

It is also important to note that the selection of the maximum number of timesteps in an episode, `max_t`, was critical for agent learning. When `max_t` originally equaled 300, the average reward plateaued at about 10. After increasing `max_t` to 20,000 and maintaining the hyperparameters listed above, the agent reached the desired average reward of +30.

2 Model Architecture

The DDPG algorithm employs two neural networks, the actor network and the critic network. For both networks, the input passes through two linear layers with relu activation. For the actor network, the hidden layers are followed by a tanh output layer. For the critic network, the hidden layers are followed by a linear output layer. The two hidden layers for each network each are comprised of 256 nodes. The code implementation is shown below.

```
def hidden_init(layer):
    fan_in = layer.weight.data.size()[0]
    lim = 1. / np.sqrt(fan_in)
    return (-lim, lim)

class Actor(nn.Module):
    """Actor (Policy) Model."""

    def __init__(self, state_size, action_size, seed, fc1_units=256, fc2_units=256):
        """Initialize parameters and build model.
```

```

Params
=====
    state_size (int): Dimension of each state
    action_size (int): Dimension of each action
    seed (int): Random seed
    fc1_units (int): Number of nodes in first hidden layer
    fc2_units (int): Number of nodes in second hidden layer
    """
    super(Actor, self).__init__()
    self.seed = torch.manual_seed(seed)
    self.fc1 = nn.Linear(state_size, fc1_units)
    self.fc2 = nn.Linear(fc1_units, fc2_units)
    self.fc3 = nn.Linear(fc2_units, action_size)
    self.reset_parameters()

def reset_parameters(self):
    self.fc1.weight.data.uniform_(*hidden_init(self.fc1))
    self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
    self.fc3.weight.data.uniform_(-3e-3, 3e-3)

def forward(self, state):
    """Build an actor (policy) network that maps states -> actions."""
    x = F.relu(self.fc1(state))
    x = F.relu(self.fc2(x))
    return F.tanh(self.fc3(x))

class Critic(nn.Module):
    """Critic (Value) Model."""

    def __init__(self, state_size, action_size, seed, fcs1_units=256, fc2_units=256):
        """Initialize parameters and build model.

        Params
        =====
            state_size (int): Dimension of each state
            action_size (int): Dimension of each action
            seed (int): Random seed
            fcs1_units (int): Number of nodes in the first hidden layer
            fc2_units (int): Number of nodes in the second hidden layer
            """
        super(Critic, self).__init__()
        self.seed = torch.manual_seed(seed)
        self.fcs1 = nn.Linear(state_size, fcs1_units)
        self.fc2 = nn.Linear(fcs1_units+action_size, fc2_units)
        self.fc3 = nn.Linear(fc2_units, 1)
        self.reset_parameters()

def reset_parameters(self):
    self.fcs1.weight.data.uniform_(*hidden_init(self.fcs1))
    self.fc2.weight.data.uniform_(*hidden_init(self.fc2))
    self.fc3.weight.data.uniform_(-3e-3, 3e-3)

def forward(self, state, action):
    """Build a critic (value) network that maps (state, action) pairs -> Q-values.
    """
    xs = F.relu(self.fcs1(state))
    x = torch.cat((xs, action), dim=1)
    x = F.relu(self.fc2(x))

```

```
return self.fc3(x)
```

3 Results

The parallel agents are able to receive an average reward (over 100 episodes, and over all 20 agents) of at least +30 in 104 episodes, as shown in Figure 1.




Figure 1: Average Score Plot

4 Future Work

The optimal value for the variable `max_t` could be explored further so that it could be reduced (thereby decreasing simulation time) while still maintaining the desired level of agent learning. From reviewing OpenAI Spinning Up’s implementation, it could also be useful to reduce the scale of the Ornstein–Uhlenbeck process noise as training advances. Spinning Up also described “Our DDPG implementation uses a trick to improve exploration at the start of training. For a fixed number of steps at the beginning, the agent takes actions which are sampled from a uniform random distribution over valid actions. After that, it returns to normal DDPG exploration [1].” This detail could also provide interesting future work.

References

- [1] Joshua Achiam. Spinning up in deep reinforcement learning. 2018.
- [2] Timothy P Lillicrap, Jonathan J Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. Continuous control with deep reinforcement learning. *arXiv preprint arXiv:1509.02971*, 2015.